

THIS IS YOUR MACHINE LEARNING SYSTEM?

YUP! YOU POUR THE DATA INTO THIS BIG
PILE OF LINEAR ALGEBRA, THEN COLLECT
THE ANSWERS ON THE OTHER SIDE.

WHAT IF THE ANSWERS ARE WRONG?

JUST STIR THE PILE UNTIL
THEY START LOOKING RIGHT.



xkcd: 1838



An Intro to BERT

**Fine-tuning Google's SOTA NLP
model to classify movie reviews**


Steven Newton
29.04.21



Making a deep contextual language model

Contextual embedding: “word embeddings are the basis of deep learning for NLP”*

Contextual word
embeddings (ELMo)



Contextual pre-training: “word masking is the crux of BERT”*

the [MASK] store called, they're running out of [MASK]

*Jacob Devlin, co-author of *BERT: Pre-training of Deep Bidirectional Transformers* (2019)

DATA

**Large Movie Review
Dataset**


ai.stanford.edu



ENVIRONMENT




python



MODELING



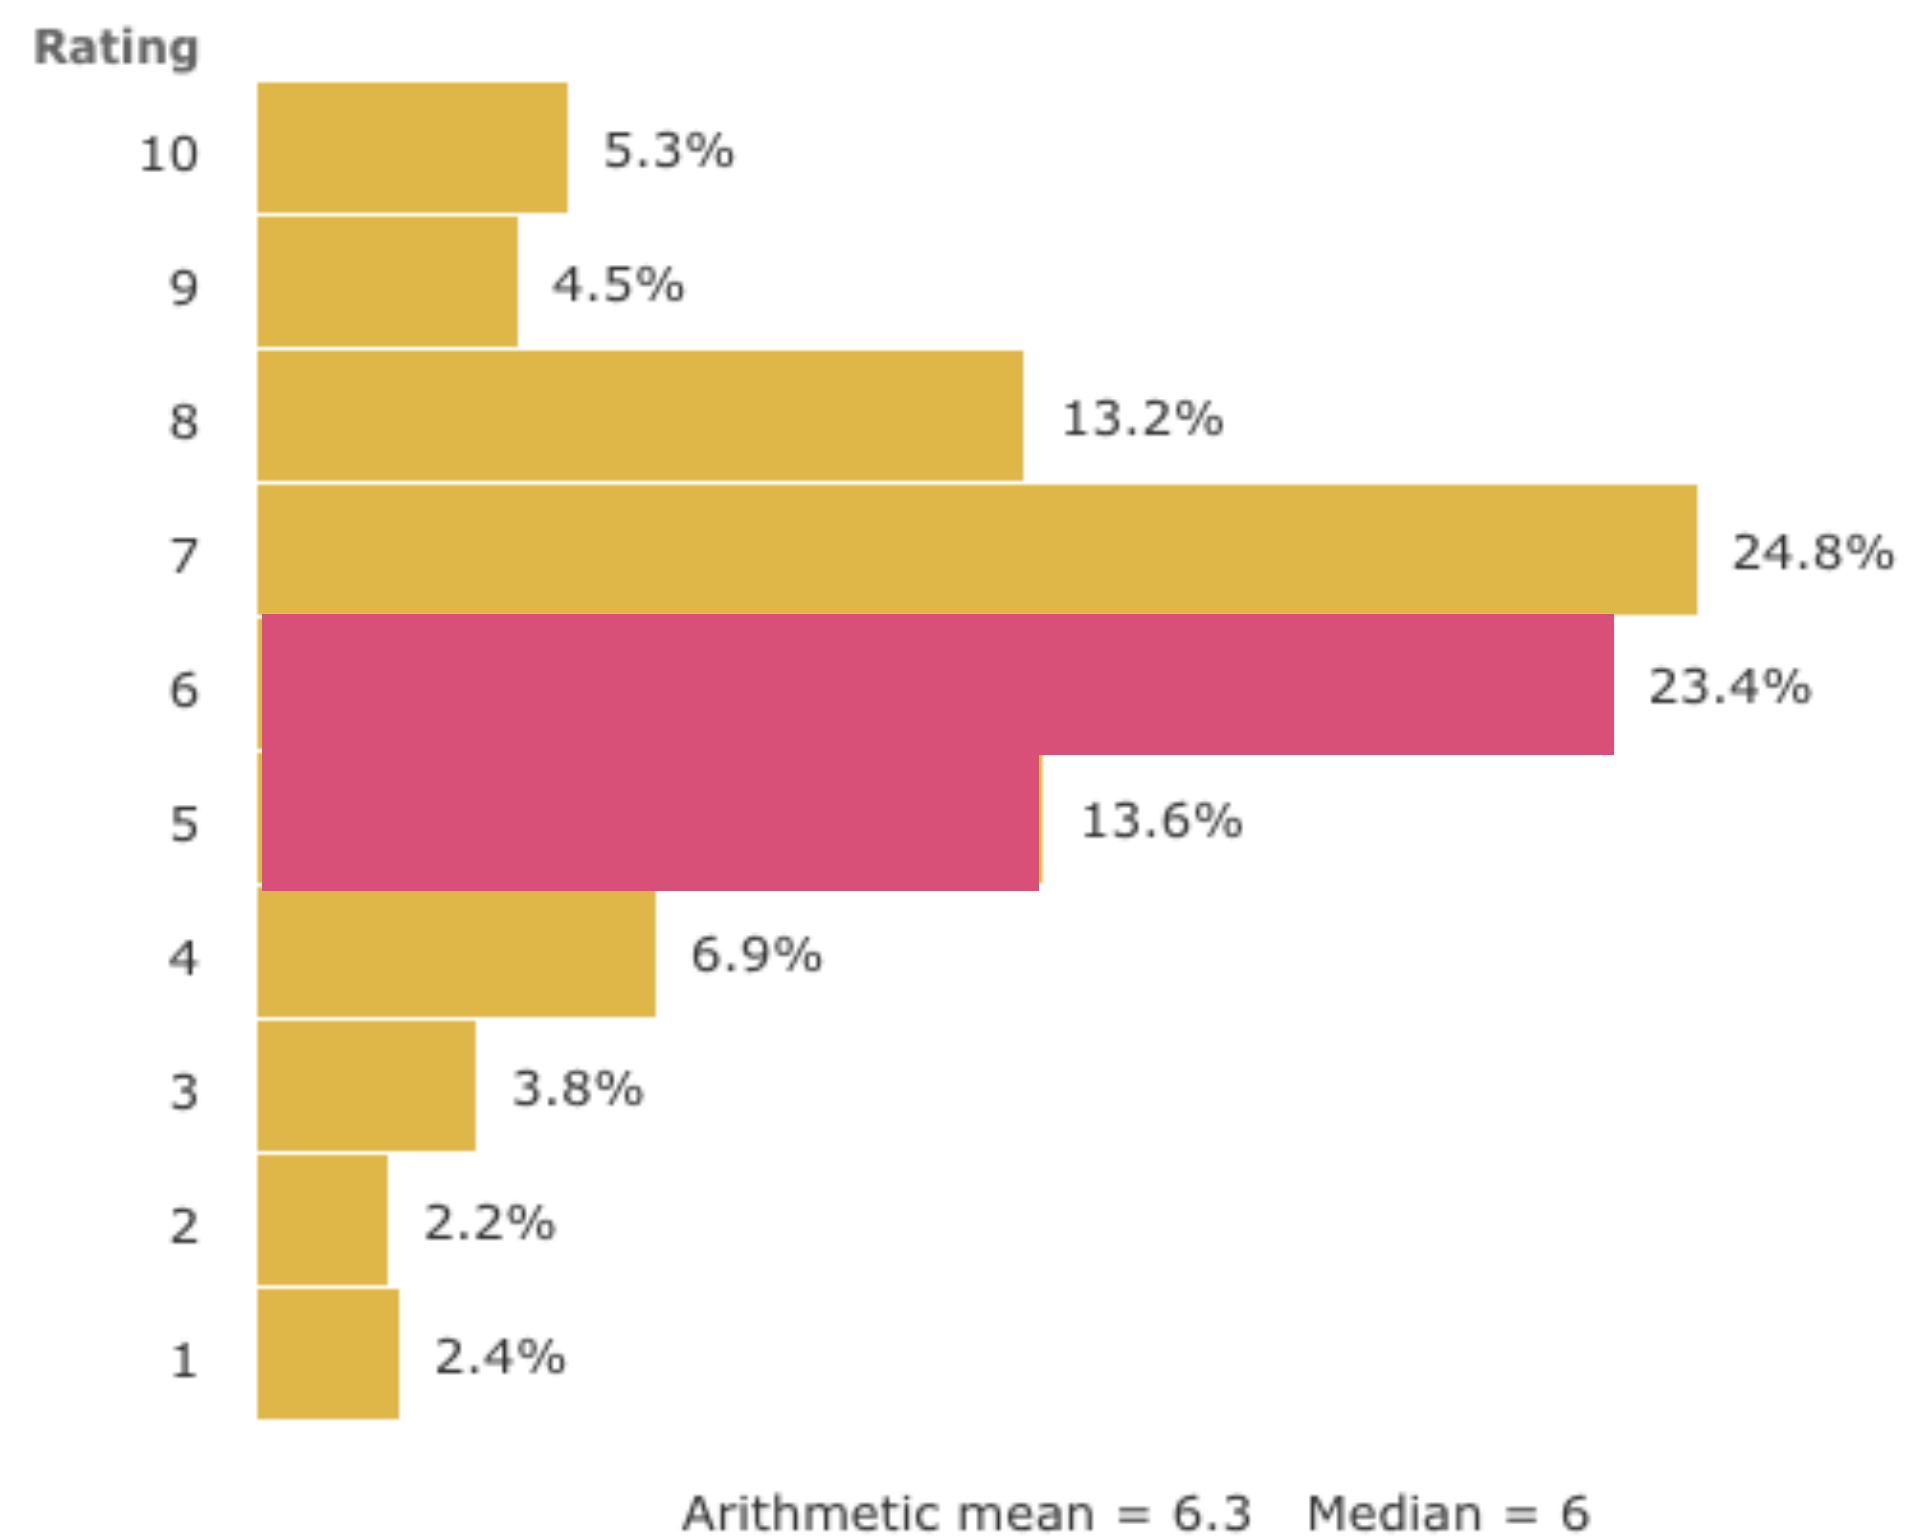


Waterworld (1995)
User Ratings
★ 6.3 ☆ Rate

IMDb Movie Reviews

IMDb Users

183,402 IMDb users have given a [weighted average](#) vote of 6.3 / 10



50,000 polar movie reviews

Binary sentiment labels

- positive
- negative

Source: Andrew L. Maas, et al. (2011).
Learning Word Vectors for Sentiment Analysis.

Results

89 percent accuracy after 20 minutes training on 1 GPU




Further gains are CTI

Ten percent of reviews truncated by 512 token input limit

- Transformer-XL
- BigBird

67%



NEGATIVE predicted as positive

model attention

Ambiguous - While the new film is gorgeous to view and the soundtrack is lovely, it is not Jane Austen's *Pride & Prejudice*...

human attention

human attention

Trolling - I only watched 15 minutes of this film...the concept alone is so original. It is so brilliant so original so GREAT. ENJOY! 3 out of 10

model attention



machines don't pick up on sarcasm



Steven Newton

`github.com/konstanzer/bert`
`steven.j36@protonmail.com`

Any sufficiently advanced technology is
indistinguishable from magic.

-Clarke's Third Law